

Problem 2: Building a Castle In The Cloud

A. Architect for the application

Please use this link and enter my code to see the diagram: <https://mermaid.live>

```
flowchart TB

%% -----
%% AWS Organization (parent)
%% -----

subgraph aws_org["AWS Organization (Multi-Account)"]

%% =====
%% Dev Account (Shortened)
%% =====

subgraph dev_acc["Dev Account"]
    direction TB
    subgraph dev_arch["Dev Core Architecture (Simplified)"]
        dev_edge["Edge (CF, WAF, ACM)"]
        dev_alb_api["ALB + API Gateway"]
        dev_eks["EKS/ECS + KEDA"]
        dev_services["Microservices & Service Mesh"]
        dev_db["Aurora/DynamoDB"]
        dev_logs["Logging (CloudWatch/X-Ray)"]
        dev_s3["S3 Data Lake (+Glue, etc.)"]
    end

    dev_end

    dev_edge --> dev_alb_api
    dev_alb_api --> dev_eks
    dev_eks --> dev_services
    dev_services --> dev_db
    dev_services --> dev_logs
    dev_s3 --> dev_logs
end

dev_end

%% =====
%% Test Account (Shortened)
```

```

%% =====
subgraph test_acc["Test Account"]
    direction TB
    subgraph test_arch["Test Core Architecture (Simplified)"]
        test_edge["Edge (CF, WAF, ACM)"]
        test_alb_api["ALB + API Gateway"]
        test_eks["EKS/ECS + KEDA"]
        test_services["Microservices & Service Mesh"]
        test_db["Aurora/DynamoDB"]
        test_logs["Logging (CloudWatch/X-Ray)"]
        test_s3["S3 Data Lake (+Glue, etc.)"]
    end

    test_edge --> test_alb_api
    test_alb_api --> test_eks
    test_eks --> test_services
    test_services --> test_db
    test_services --> test_logs
    test_s3 --> test_logs
end

%% =====
%% Prod Account (Full Detail)
%% =====
subgraph prod_acc["Prod Account"]
    direction TB

    %% -----
    %% Subgraph: Edge
    %% -----
    subgraph prod_edge["Edge Layer"]
        cf["Amazon CloudFront"]
        waf["AWS WAF"]
        acm["AWS Certificate Manager"]
        route53["Route 53 (DNS + Health Checks)"]
    end

    route53 --> cf
    cf --> waf
    waf --> acm

    %% -----

```

```

%% Subgraph: API & Ingress
%% -----
subgraph prod_api["API & Ingress"]
    alb["Application Load Balancer"]
    apigw["Amazon API Gateway"]
    cognito["AWS Cognito / SSL Termination"]
end

acm --> alb
alb --> apigw
apigw --> cognito

%% -----
%% Subgraph: Microservices Orchestration
%% -----
subgraph prod_micro["Microservices Orchestration"]
    eks["EKS (or ECS) + KEDA"]
    svcMesh["Service Mesh (Istio/App Mesh)"]
    intApi["Internal API Gateway <br/> for Microservices"]
end

cognito --> eks
eks --> svcMesh
svcMesh --> intApi

%% -----
%% Subgraph: Core Trading Components
%% -----
subgraph prod_trading["Core Trading Components"]
    matchEng["Real-Time Matching Engine <br/>(Dedicated EC2 in
ASG) "]
    llm["Low Latency Messaging <br/>(Kinesis / Kafka)"]
end

eks --> llm
llm --> matchEng

%% -----
%% Subgraph: Cache & Session
%% -----
subgraph prod_cache["Cache & Session Management"]
    redis["Amazon ElastiCache (Redis)"]
end

```

```

intApi --> redis
matchEng --> redis

%% -----
%% Subgraph: Data & Transactional
%% -----
subgraph prod_data["Data & Transactional Storage"]
    aurora["Aurora / DynamoDB (Sharded/Global)"]
    acidDB["ACID-compliant Order DB"]
end

intApi --> aurora
matchEng --> aurora

%% -----
%% Subgraph: Monitoring & Security
%% -----
subgraph prod_sec["Security & Audit"]
    logs["CloudWatch, X-Ray"]
    enc["Encryption & IAM"]
end

eks --> logs
matchEng --> logs
aurora --> logs

%% -----
%% Subgraph: Historical & Analytics
%% -----
subgraph prod_analytics["Historical & Analytics"]
    s3["S3 w/ Lifecycle Policies"]
    glue["AWS Glue (Catalog/ETL)"]
    redshift["Amazon Redshift"]
    athena["Athena (Ad-hoc Queries)"]
    kda["Kinesis Data Analytics"]
end

s3 --> glue
s3 --> athena
glue --> redshift
kda --> s3

```

```

    %% Full Connections Inside Prod
    acm --> alb
    alb --> apigw
    apigw --> cognito
    cognito --> eks
    eks --> svcMesh
    svcMesh --> intApi
    intApi --> aurora
    intApi --> redis
    llm --> matchEng
    matchEng --> redis
    matchEng --> aurora
    aurora --> logs
    redis --> logs

    %% Additional flows to analytics
    intApi --> s3
    matchEng --> s3
    aurora --> s3
    s3 --> logs
end

%% Inter-Account Networking
tgw["Transit Gateway or VPC Peering"]
dev_acc --> tgw
test_acc --> tgw
prod_acc --> tgw
end

%% -----
%% Security & Governance (Central)
%% -----
subgraph sec_gov["Security & Governance"]
    secHub["AWS Security Hub"]
    guardDuty["Amazon GuardDuty"]
    awsConfig["AWS Config"]
    iamScp["IAM Roles & SCPs (Organizations)"]
end

aws_org --> sec_gov

```

1. Multi-Account Structure

In a mature AWS setup, you often separate environments (e.g., Dev, Test, Prod) into **different AWS accounts** to:

- **Isolate** resources so that issues or changes in Dev/Test do not affect Production.
- **Improve security** by limiting cross-environment access.
- **Simplify billing** (separate cost tracking for each environment).
- **Enforce compliance** using organizational policies.

Dev Account

- **Purpose:** Used for active development, experimentation, and integration testing. Developers have more freedom here to spin up services, test features, and iterate quickly.

Test Account

- **Purpose:** Dedicated for QA, staging, and performance testing. Ensures changes are validated (functional, performance, security) before pushing to production.

Prod Account

- **Purpose:** Hosts your live, customer-facing or mission-critical workloads. Must be highly secured, monitored, and follow strict change-management processes to ensure uptime and data integrity.

Transit Gateway / VPC Peering

- **Purpose:** Provides **network connectivity** between the separate AWS accounts.
 - **Transit Gateway** scales better for multiple VPCs; **VPC Peering** is simpler for one-to-one connections.
 - **Benefit:** Allows shared services or data replication between environments.
-

2. Edge Layer

Route 53 (DNS & Health Checks)

- **Purpose:** AWS' DNS service to direct user traffic to your application endpoints.
- **Key Capabilities:**
 - **Health Checks:** Automatically check endpoint health and reroute if one endpoint fails.
 - **Routing Policies:** Simple, Weighted, Geolocation, Latency-based, Failover, etc., to optimize domain traffic management.

Amazon CloudFront

- **Purpose:** Content Delivery Network (CDN) that caches and delivers content (static or dynamic) from edge locations close to users.
- **Benefit:** Improves **performance** (reduced latency) and **reduces load** on origins (e.g., ALB, S3).

AWS WAF (Web Application Firewall)

- **Purpose:** Protects web applications from common attacks (SQL injection, XSS, bots, etc.).
- **Integration:** Typically attaches to CloudFront or Application Load Balancers for centralized inspection of HTTP/S traffic.

AWS Certificate Manager (ACM)

- **Purpose:** Issues and manages **SSL/TLS certificates** for your domains.
 - **Benefit:** Automates certificate **renewal**, integrates seamlessly with AWS services like CloudFront, ALB, and API Gateway.
-

3. API & Ingress

Application Load Balancer (ALB)

- **Purpose:** Distributes incoming HTTP(S) requests across multiple targets (EC2, ECS, EKS pods).
- **Key Features:** Path-based routing, host-based routing, WebSocket support, integration with AWS WAF.

Amazon API Gateway

- **Purpose:** Manages **REST, HTTP, or WebSocket** APIs. Offers features like **throttling, request validation, and usage plans**.
- **Benefit:** Simplifies exposing microservices as well-defined APIs; can integrate with **AWS Lambda, ECS/EKS, or VPC Link** to backend services.

AWS Cognito

- **Purpose:** Provides **user authentication, authorization, and user management** (signup/signin flows, secure tokens).
 - **Integration:** Often used with API Gateway to handle **JWT-based** auth or federated SSO (e.g., Social logins, SAML).
-

4. Microservices Orchestration

Amazon EKS (or ECS) + KEDA

- **Purpose:**
 - **EKS:** Managed Kubernetes control plane, schedules containerized apps (pods) across EC2 or Fargate.
 - **ECS:** Alternative AWS container service with simpler operational overhead.
 - **KEDA** (Kubernetes-based Event Driven Autoscaling): Monitors external event sources (e.g., SQS, Kinesis) and **auto-scales** pods based on queue depth or event volume.
- **Benefit:** Scalability, flexibility, standardized container deployment.

Service Mesh (Istio / App Mesh)

- **Purpose:** Controls **service-to-service communication** with features like **traffic routing**, **mTLS encryption**, **observability** (metrics, logs, tracing), and **security policies**.
- **Benefit:** Offloads complex networking logic from individual microservices to a **sidecar proxy**.

Internal API Gateway

- **Purpose:** Some organizations prefer an **internal** gateway for microservices to unify internal API routes, apply consistent policies, or handle zero-trust networking behind the service mesh.
 - **Note:** Depending on complexity, the service mesh alone might handle these responsibilities without a separate gateway.
-

5. Core Trading Components (Example: High-performance, finance-related)

Real-Time Matching Engine (EC2 in ASG)

- **Purpose:** Processes **trade orders** and matches them in real-time with minimal latency.
- **Why EC2:** If you need **low-level OS access**, specialized libraries, or consistent ultra-low latency, a dedicated EC2 cluster may outperform containers.
- **Auto Scaling Group (ASG):** Automatically adds or removes EC2 instances based on CPU, memory, or custom metrics (e.g., orders/sec).

Low Latency Messaging (Kinesis / Kafka)

- **Purpose:** Ingestion and distribution of **streaming data** (trades, events) at scale.

- **Kinesis** is a fully managed AWS service; **Apache Kafka** (or MSK) offers a popular open-source alternative for advanced streaming.
 - **Benefit:** Real-time data pipelines, decoupled microservices, backpressure handling, replay capabilities.
-

6. Cache & Session Management

Amazon ElastiCache (Redis)

- **Purpose:** High-performance **in-memory** caching layer for session data, frequently accessed queries, or real-time analytics.
 - **Benefit:** Reduces latency, offloads reads from the database, can manage user sessions at scale.
-

7. Data & Transactional Storage

Aurora / DynamoDB (Sharded/Global)

- **Aurora** (Relational): MySQL/PostgreSQL-compatible. Offers **ACID** transactions, read replicas, and global database features for **low-latency cross-region** reads and rapid failover.
- **DynamoDB** (NoSQL): Scales horizontally for key-value or document store workloads, supports global tables for multi-region replication, and provides **eventual** or **strong** consistency.
- **Sharding:** A strategy to split large volumes of data across multiple DB instances or tables.

ACID-Compliant Order DB

- **Purpose:** If you need strict, **atomic** transaction guarantees (e.g., for financial trades), an **ACID** database ensures data consistency even in high-throughput scenarios.
 - **Benefit:** Preserves transactional integrity (e.g., no partial trades in the system).
-

8. Security & Audit

CloudWatch, X-Ray

- **Purpose:** Centralized **logging**, **metrics**, and **tracing** for distributed systems.
- **CloudWatch:** Collects logs (e.g., from ECS tasks, EKS pods), infrastructure metrics, and custom application metrics.

- **X-Ray:** Enables **distributed tracing** for microservices to diagnose performance bottlenecks and trace requests end-to-end.

Encryption & IAM

- **Purpose:**
 - **Encryption:** Secure data at rest (e.g., in EBS, RDS, S3) and in transit (TLS).
 - **IAM:** Access control for users, services, and resources with **least-privilege** policies.
 - **Benefit:** Helps meet compliance (HIPAA, PCI-DSS, etc.) and ensures only authorized actors can access data.
-

9. Historical & Analytics

Amazon S3 (w/ Lifecycle Policies)

- **Purpose:** Durable, scalable object storage for raw data, logs, archives, backups.
- **Lifecycle Policies:** Automatically transition older data to cheaper classes (e.g., Glacier).

AWS Glue (Catalog/ETL)

- **Purpose:** Data **catalog** and ETL service. Glue **crawlers** detect schema in S3 data, making it queryable in Athena or Redshift Spectrum.
- **Benefit:** Simplifies building a **data lake** in S3.

Amazon Redshift

- **Purpose:** Fully managed, **petabyte-scale** data warehouse for analytics and BI dashboards.
- **Integration:** Can query data in S3 via **Redshift Spectrum**.

Amazon Athena

- **Purpose:** Serverless, interactive query service for data in S3.
- **Benefit:** Great for ad-hoc queries, does not require managing any infrastructure.

Kinesis Data Analytics

- **Purpose:** Real-time analytics on streaming data (from Kinesis or Kafka).
 - **Benefit:** Perform **windowed** operations, real-time metrics, or anomaly detection without building a complex pipeline.
-

10. Security & Governance (Organization-wide)

AWS Security Hub

- **Purpose:** A unified dashboard that aggregates security findings from **GuardDuty**, **Inspector**, **Macie**, and other tools. Provides a **single pane of glass** for compliance checks (CIS, PCI, etc.).

Amazon GuardDuty

- **Purpose:** Continuous **threat detection** service that uses machine learning to identify malicious activity or unauthorized behavior.
- **Monitors** CloudTrail, VPC Flow Logs, DNS logs, etc.

AWS Config

- **Purpose:** Tracks **configuration changes** of AWS resources. Allows you to define rules (e.g., "S3 buckets must not be public") to enforce compliance.
- **Benefit:** Provides a historical record of resource config changes for audits.

IAM Roles & Service Control Policies (SCPs)

1. **Purpose:**
 - a. **IAM Roles:** Grants permissions to AWS services or users based on the principle of **least privilege**.
 - b. **SCPs:** Organizational-level "guardrails" that limit what actions are allowed across **all** accounts or within specific Organizational Units (OUs).
2. **Benefit:** Prevents accidental or malicious usage of restricted services/actions across Dev/Test/Prod.

B. The 5 main features that we can cover with this architect

1. Real-Time Order Matching and Execution

Key Services

- **Real-Time Matching Engine (EC2 in ASG)**
- **Low-Latency Messaging (Kinesis/Kafka)**
- **Microservices Orchestration (EKS/ECS + Service Mesh)**
- **ElastiCache (Redis)**

Explanation

1. **Dedicated Matching Engine:**
 - Placed on **EC2 instances in an Auto Scaling Group**, allowing you to run specialized low-latency code (e.g., C++/Rust) without container overhead.
 - You can fine-tune OS settings, CPU pinning, and networking (ENA drivers) to minimize jitter—essential for high-frequency trading.
2. **Low-Latency Messaging:**
 - **Kinesis** or **Kafka** ensures **real-time ingestion** of new orders.
 - Messaging decouples order placement from the matching engine, preventing direct coupling that can lead to bottlenecks.
 - **Partitioning** or **shards** can align with specific trading pairs or symbol sets, keeping throughput high and latency predictable.
3. **Microservices & Orchestration:**
 - **EKS/ECS** handles ancillary microservices (e.g., user profile, risk checks, trade history) that interact with the matching engine via APIs or messaging queues.
 - A **Service Mesh** like Istio or AWS App Mesh ensures consistent, **secure** service-to-service communication with built-in **observability** (metrics/traces).
4. **In-Memory Cache (Redis):**
 - Used for **hot data** (e.g., open orders, rapidly updated user balances).
 - Reduces round trips to the database, keeping latencies low for real-time lookups.

Result: By decoupling order ingestion (Kinesis/Kafka) from order matching (EC2/ASG) and caching (Redis), you get **very fast, real-time** execution with minimal overhead.

2. Transaction Integrity and ACID Compliance

Key Services

- Amazon Aurora (SQL) / DynamoDB (NoSQL with Transactions)
- ACID-Compliant Order & Transaction Database
- Microservices Orchestration

Explanation

1. **ACID-Compliant Database:**
 - **Aurora** (MySQL/Postgres-compatible) provides **strong ACID guarantees** with built-in replication, point-in-time recovery, and multi-AZ durability.
 - If using **DynamoDB**, transactions at the item-level can **atomically** update multiple items, ensuring partial updates do not occur.
2. **Atomic Operations:**
 - In a trading context, partial trades or inconsistent states can cause major discrepancies.
 - Aurora's **row-level locking** or DynamoDB's **transaction support** helps maintain data integrity (e.g., account balance updates, order status changes).
3. **Microservices Boundaries:**
 - Each microservice (e.g., "Order Service," "Portfolio Service," etc.) interacts with a transactional store.
 - Patterns like **sagas** or **two-phase commit** (if needed) coordinate multi-service transactions—though often you keep each transaction bounded in one domain service to keep it simpler.
4. **Event Sourcing & Auditing:**
 - Data changes can be **streamed** to Kinesis/Firehose for an immutable record of all trades.
 - This also helps with compliance/regulatory audits.

Result: Ensuring each trade or balance update is handled in a **transactional manner** prevents double-spend errors and ensures **full data consistency** for critical financial operations.

3. Handling High Traffic and Peak Loads

Key Services

- Autoscaling (EKS/ECS with KEDA / EC2 ASG)
- Edge (CloudFront, WAF, ACM) + ALB/API Gateway
- ElastiCache
- Kinesis/Kafka (Partitioned Streams)

Explanation

1. **Autoscaling Everywhere:**

- **EKS/ECS** can horizontally scale microservices based on CPU, memory, or custom metrics (via **KEDA**).
 - **EC2 Auto Scaling Groups** add or remove matching engine nodes under heavy workloads (e.g., major market movements).
 - **Kinesis** or **Kafka** partition scaling ensures message throughput can handle high volumes of orders.
2. **Edge Layer:**
- **CloudFront** caches static and semi-static content (e.g., the trading dashboard's static assets), reducing load on application backends.
 - **AWS WAF** filters malicious requests early, preventing resource waste from potential DDoS attempts or bot traffic.
3. **Caching:**
- **Redis** for frequently accessed data (e.g., market data, real-time order book snapshots).
 - Minimizes repeated queries on primary databases, thus **reducing latency** and improving throughput under surge traffic.
4. **API Gateway + ALB:**
- Splits traffic: high-throughput, lower-level requests might route through an **ALB** directly, while advanced API-based flows go through **API Gateway**.
 - **Usage plans** and **throttling** help prevent service overload.

Result: With robust **scaling** at every layer and effective **caching/streaming**, the system can handle massive peaks—akin to what major exchanges face during volatile crypto market events.

4. Advanced Security and DDoS Mitigation

Key Services

- **AWS WAF + CloudFront**
- **AWS Shield (Standard or Advanced)**
- **IAM Roles & SCPs**
- **Security Hub, GuardDuty, Config**
- **Encryption (ACM, TLS, KMS)**

Explanation

1. **WAF + CloudFront:**
 - **WAF** can block malicious patterns (SQL injection, XSS) at the edge before they reach your application.
 - **CloudFront** offloads TLS termination, absorbs spikes, and integrates with **AWS Shield** to mitigate DDoS at the network layer (L3/L4).
2. **Advanced Threat Detection:**
 - **GuardDuty** continuously analyzes VPC flow logs, CloudTrail logs, and DNS queries for suspicious activities (e.g., port scans, anomalous traffic).

- Findings feed into **Security Hub**, giving a unified security dashboard for all AWS accounts.
- 3. **IAM & SCPs:**
 - **IAM** roles with least-privilege policies ensure microservices and DevOps engineers have only the necessary access.
 - **Service Control Policies (SCPs)** restrict high-risk actions or services across Dev/Test/Prod at the organizational level.
- 4. **Encryption:**
 - Data at rest in **Aurora**, **DynamoDB**, **S3**, and **ElastiCache** can be encrypted with AWS KMS.
 - In-transit encryption (TLS/SSL) enforced at CloudFront, ALB, or API Gateway using **AWS Certificate Manager**.
 - **mTLS** inside the service mesh can protect internal service-to-service traffic from snooping.
- 5. **DDoS Mitigation:**
 - **AWS Shield Standard** automatically included with CloudFront and Route 53 to handle common DDoS attacks.
 - **AWS Shield Advanced** (optional) provides more sophisticated detection and cost protection during large-scale attacks.

Result: An advanced multi-layer security approach from edge to data ensures the platform remains **secure and compliant**, even under targeted DDoS or other sophisticated threats.

5. Crypto-Specific Features and Compliance

Key Services

- **Multi-Account Setup** (Dev/Test/Prod)
- **Audit Logging** (CloudWatch, X-Ray, S3)
- **IAM, Config, KMS**
- **Data Lake** (S3, Glue, Redshift) for analytics and forensics
- **(Optional) HSM / AWS CloudHSM** for private key management if needed

Explanation

1. **Regulatory Compliance & Audits:**
 - Many crypto exchanges must adhere to KYC/AML regulations, **FATF guidelines**, or local financial authorities.
 - You can keep user data, logs, and transaction records in **separate accounts** (lower blast radius), track changes with **Config**, and gather logs in **CloudWatch** or **S3** for audit trails.
2. **Data Lake for Transaction & Trade History:**
 - Storing all trades and user events in **S3** (with **Glue Data Catalog**) provides an **immutable** record of historical data for compliance and forensic analysis.
 - **Encryption at rest** ensures customer data confidentiality.

- **Lifecycle Policies** automatically transition older data to cheaper storage (Glacier), meeting **long-term** record-keeping requirements.
- 3. **Secure Key Management:**
 - **AWS KMS** can manage encryption keys for data at rest.
 - If you need hardware-level cryptographic modules for cold or warm wallets, **AWS CloudHSM** or external HSMs can integrate with your exchange.
 - This is critical for **custody of crypto assets**, though many large exchanges use separate hardware wallets or third-party custody solutions.
- 4. **Isolated Dev/Test/Prod:**
 - Ensures **production** cryptographic keys, wallets, and user data remain **completely separate** from development environments, reducing the risk of internal fraud or accidental data leaks.
- 5. **Integration with Identity Providers:**
 - **Cognito** can integrate with user pool logic for KYC checks or 2FA (MFA), while additional checks or AML logic can be implemented at the application layer or via 3rd-party APIs.

Result: By leveraging **multi-account isolation**, **secure data storage**, **robust logging**, and optional **CloudHSM** for key management, your architecture meets the **unique compliance and security needs** of a crypto trading platform.