

# Problem 3: Diagnose Me Doctor

Below is a quick-reference grid for the two most common causes:

Root Cause	Impacts	Recovery Actions
Misconfiguration / Memory Leak	Gradual climb to 99% memory, OOM kills, dropped requests.	<ul style="list-style-type: none"><li>- Restart NGINX</li><li>- Adjust or remove problematic modules</li><li>- Tune workers, caching</li><li>- Patch/upgrade NGINX</li></ul>
Excessive / Malicious Traffic	Rapid or sustained spike in memory usage, request timeouts, and potential downtime.	<ul style="list-style-type: none"><li>- Block or throttle suspicious IPs</li><li>- Scale horizontally (more load balancers)</li><li>- Harden configuration with rate-limiting</li><li>- Use DDoS mitigation service</li></ul>

## 1. Verify and Quantify the Issue

### 1. Check Monitoring Dashboards

- Confirm that memory usage is indeed at 99% and not a false alarm.
- Compare current usage with historical trends to see if this is a one-time spike or a recurring issue.

### 2. Inspect Memory on the VM

- Run `free -h` or `vmstat` to verify total memory, used memory, and any swap usage.
- Use `top` / `htop` to see real-time memory consumption by process.

### 3. Check Processes Sorted by Memory

- `ps aux --sort=-%mem | head -n 10` to identify the top memory-consuming processes.
- Confirm that NGINX (or one of its worker processes) is the culprit.

---

## 2. Gather Logs and System Information

### 1. System Logs

- Review `/var/log/syslog` or `journalctl -xe` for any OOM (Out-Of-Memory) kill events or errors pointing to resource exhaustion.
- Look at `dmesg` for kernel messages that might show memory allocation or OOM kill details.

## 2. NGINX Logs

- Check the NGINX error log (commonly `/var/log/nginx/error.log`) for memory allocation issues, repeated “worker process exited” messages or warnings about timeouts/overloads.
- Examine the access log (`/var/log/nginx/access.log`) for unusual spikes in requests, and patterns that might indicate a flood of traffic or DDoS attempts.

## 3. Configuration Files

- Open `/etc/nginx/nginx.conf` to review the main configuration:
    - `worker_processes`, `worker_connections`, and any module-specific settings.
  - Look at additional config files in `/etc/nginx/conf.d` or `/etc/nginx/sites-enabled` for caching, Lua modules, or other features that might consume extra memory.
- 

# 3. Narrow Down the Root Cause

## 3.1 Confirm Whether It's a Misconfiguration or Memory Leak

### 1. Inspect NGINX Worker Processes

- Run `ps aux | grep nginx` to see each worker process's memory usage.
- If one or more workers are continuously growing in memory usage, suspect a memory leak or misconfiguration (e.g., excessive caching).

### 2. Audit Custom Modules

- Check if you're using any third-party modules (e.g., Lua, Redis, or advanced caching).
- Temporarily disable or revert to a vanilla NGINX configuration to see if the memory usage improves.

### 3. Tweak NGINX Worker Settings

- If `worker_processes` is manually set too high (e.g., 16 workers on a small VM), reduce it to align with available CPU cores or use `worker_processes auto`;
- Review `worker_connections` to ensure it isn't excessively high for your traffic patterns.

### 4. Look for Known Bugs

- Check your NGINX version and any modules against known issues in official or community forums.
- Update or patch if there's a reported memory leak in a specific version or module.

## Potential Outcomes

- **Misconfiguration:** Memory usage drops after correcting `worker_processes` or disabling a problematic module.

- **Memory Leak in a Module:** Memory usage continues to climb over time. A software upgrade or patch may be required.
- 

## 3.2 Check for Excessive or Malicious Traffic

### 1. Traffic Analysis

- Inspect access logs to identify a sudden surge in traffic or repetitive requests from the same IP ranges.
- Use `netstat -antp` or `ss -tunap` to see the number of active connections and their states (ESTABLISHED, TIME\_WAIT, etc.).

### 2. Rate of Requests

- If your normal traffic is, say, 1K requests/min, but logs show 5x that amount, suspect a DDoS or heavy legitimate spike (e.g., a marketing campaign).
- Compare request patterns with standard load to confirm abnormal spikes.

### 3. Check Upstream Errors

- If you have health checks or an upstream service that is failing, NGINX might be retrying connections excessively, leading to ballooning memory usage.

## Potential Outcomes

- **Legitimate Traffic Spike:** High usage caused by real demand (e.g., a product launch).
  - **DDoS or Malicious Traffic:** Abnormal IP patterns, large numbers of small requests, or suspicious user agents flooding the service.
- 

## 4. Immediate Remediation Steps

### 1. Restart NGINX

- `sudo systemctl restart nginx` or `sudo service nginx restart` can temporarily free memory used by worker processes.
- Monitor to see if memory usage quickly ramps up again (indicating a persistent leak or recurring spike in traffic).

### 2. Enable or Increase Swap (Short-Term Relief)

- If you're running critically low on RAM, adding swap (`sudo fallocate -l 2G /swapfile`, then `sudo mkswap /swapfile && sudo swapon /swapfile`) can prevent immediate OOM kills.
- Note that swap is slower than RAM—this is a band-aid, not a permanent solution.

### 3. Block or Rate-Limit Malicious IPs

- If logs indicate malicious traffic from specific IP ranges, block them temporarily with firewall rules (e.g., `iptables` or your cloud's security groups).

- Consider using `limit_req_zone` or `limit_conn_zone` in NGINX to throttle suspicious traffic.
4. **Scale Up Temporarily**
- If you're on a cloud platform, increase the VM size (RAM) or add more load balancer nodes to handle the surge until you can implement a more permanent fix.
- 

## 5. Long-Term Fixes

1. **Optimize NGINX Configuration**
  - Right-size `worker_processes` and `worker_connections`.
  - Tune keep-alive settings (`keepalive_timeout`) to ensure connections aren't held open unnecessarily.
  - If caching, configure cache size limits and eviction policies carefully.
2. **Upgrade or Patch NGINX / Modules**
  - If a memory leak is identified in a specific version, upgrade to the latest stable release.
  - Remove or replace buggy third-party modules.
3. **Set Up Autoscaling**
  - Use a managed load balancer service or create an autoscaling group that spins up additional instances when traffic spikes.
  - Implement a DNS-based round-robin or cloud load-balancing solution for better high-availability.
4. **Harden Against DDoS**
  - Enable DDoS protection at the cloud or CDN level (e.g., AWS Shield, Cloudflare, etc.).
  - Configure rate-limiting rules, IP reputation checks, and WAF (Web Application Firewall) to filter malicious traffic.
5. **Improve Observability**
  - Implement metrics collection and analysis (e.g., Prometheus + Grafana).
  - Create alerts on memory usage, connection counts, 5xx error rates, etc., to catch issues early.